

Analysis of Yelp Business Intelligence Data

In this project, I analyzed a subset of the Yelp's business, reviews and user data.

The dataset originally comes from [Kaggle](#) and it has been uploaded into an S3 bucket:

```
s3://yelpreviewdataset/yelp_academic_dataset_business.json\
s3://yelpreviewdataset/yelp_academic_dataset_review.json\
s3://yelpreviewdataset/yelp_academic_dataset_user.json
```

Part I: Installation and Initial Setup

1. Install Packages

```
In [1]: from pyspark.sql import SparkSession
spark = SparkSession \
    .builder \
    .appName("Analysis of Yelp Business") \
    .config("spark.some.config.option", "some-value") \
    .getOrCreate()
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
sc.install_pypi_package("seaborn==0.10.0")
sc.list_packages()
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
0	application_1606187319027_0001	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Collecting pandas==1.0.3

Downloading https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl (10.0MB)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)

Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)

Downloading https://files.pythonhosted.org/packages/d4/70/d60450c3dd48ef87586924207ae8907090de0b306af2bce5d134d78615cb/python_dateutil-2.8.1-py2.py3-none-any.whl (227kB)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)

Installing collected packages: python-dateutil, pandas

Successfully installed pandas-1.0.3 python-dateutil-2.8.1

Collecting matplotlib==3.2.1

Downloading https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl (12.4 MB)

Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from matplotlib==3.2.1)

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Downloading <https://files.pythonhosted.org/packages/8a/bb/488841f56197b13700afd5658fc279a2025a39e22449b7cf29864669b15d/pyparsing-2.4.7-py2.py3-none-any.whl> (67kB)

Collecting cycler>=0.10 (from matplotlib==3.2.1)
 Downloading <https://files.pythonhosted.org/packages/f7/d2/e07d3ebbd7af696440ce7e754c59dd546fffe1bbe732c8ab68b9c834e61/cycler-0.10.0-py2.py3-none-any.whl>
 Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)
 Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)
 Downloading https://files.pythonhosted.org/packages/d2/46/231de802ade4225b76b96cffe419cf3ce52bbe92e3b092cf12db7d11c207/kiwisolver-1.3.1-cp37-cp37m-manylinux1_x86_64.whl (1.1MB)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)
 Installing collected packages: pyparsing, cycler, kiwisolver, matplotlib
 Successfully installed cycler-0.10.0 kiwisolver-1.3.1 matplotlib-3.2.1 pyparsing-2.4.7

Collecting seaborn==0.10.0
 Downloading <https://files.pythonhosted.org/packages/70/bd/5e6bf595fe6ee0f257ae49336dd180768c1ed3d7c7155b2fdf894c1c808a/seaborn-0.10.0-py3-none-any.whl> (215kB)
 Requirement already satisfied: pandas>=0.22.0 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from seaborn==0.10.0)
 Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.10.0)
 Collecting scipy>=1.0.1 (from seaborn==0.10.0)
 Downloading https://files.pythonhosted.org/packages/dc/7e/8f6a79b102ca1ea928bae8998b05bf5dc24a90571db13cd119f275ba6252/scipy-1.5.4-cp37-cp37m-manylinux1_x86_64.whl (25.9MB)
 Requirement already satisfied: matplotlib>=2.1.2 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from seaborn==0.10.0)
 Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)
 Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: cycler>=0.10 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1606188495504-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas>=0.22.0->seaborn==0.10.0)
 Installing collected packages: scipy, seaborn
 Successfully installed scipy-1.5.4 seaborn-0.10.0

Package	Version
-----	-----
beautifulsoup4	4.9.1
boto	2.49.0
click	7.1.2
cycler	0.10.0
jmespath	0.10.0
joblib	0.16.0
kiwisolver	1.3.1
lxml	4.5.2
matplotlib	3.2.1
mysqlclient	1.4.2
nlTK	3.5
nose	1.3.4
numpy	1.16.5
pandas	1.0.3
pip	9.0.1
py-dateutil	2.2
pyparsing	2.4.7
python-dateutil	2.8.1
python37-sagemaker-pyspark	1.4.0
pytz	2020.1
PyYAML	5.3.1
regex	2020.7.14

scipy	1.5.4
seaborn	0.10.0
setuptools	28.8.0
six	1.13.0
soupsieve	1.9.5
tqdm	4.48.2
wheel	0.29.0
windmill	1.6

2. Importing

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pyspark.sql.functions as F
from scipy import stats
from scipy.stats import norm, skew
from pyspark.sql.functions import explode, split, desc, col, avg, udf, when
from pyspark.sql.types import IntegerType, StringType, DoubleType
```

3. Loading Data

```
In [3]: business_df = spark.read.json('s3://yelpreviewdataset/yelp_academic_dataset_business.js
```

4. Overview of Data

```
In [4]: print(f'Columns: {len(business_df.columns)} | Rows: {business_df.count():,}')
```

Columns: 14 | Rows: 209,393

```
In [5]: business_df.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
```

```

|-- GoodForDancing: string (nullable = true)
|-- GoodForKids: string (nullable = true)
|-- GoodForMeal: string (nullable = true)
|-- HairSpecializesIn: string (nullable = true)
|-- HappyHour: string (nullable = true)
|-- HasTV: string (nullable = true)
|-- Music: string (nullable = true)
|-- NoiseLevel: string (nullable = true)
|-- Open24Hours: string (nullable = true)
|-- OutdoorSeating: string (nullable = true)
|-- RestaurantsAttire: string (nullable = true)
|-- RestaurantsCounterService: string (nullable = true)
|-- RestaurantsDelivery: string (nullable = true)
|-- RestaurantsGoodForGroups: string (nullable = true)
|-- RestaurantsPriceRange2: string (nullable = true)
|-- RestaurantsReservations: string (nullable = true)
|-- RestaurantsTableService: string (nullable = true)
|-- RestaurantsTakeOut: string (nullable = true)
|-- Smoking: string (nullable = true)
|-- WheelchairAccessible: string (nullable = true)
|-- WiFi: string (nullable = true)
-- business_id: string (nullable = true)
-- categories: string (nullable = true)
-- city: string (nullable = true)
-- hours: struct (nullable = true)
|   |-- Friday: string (nullable = true)
|   |-- Monday: string (nullable = true)
|   |-- Saturday: string (nullable = true)
|   |-- Sunday: string (nullable = true)
|   |-- Thursday: string (nullable = true)
|   |-- Tuesday: string (nullable = true)
|   |-- Wednesday: string (nullable = true)
-- is_open: long (nullable = true)
-- latitude: double (nullable = true)
-- longitude: double (nullable = true)
-- name: string (nullable = true)
-- postal_code: string (nullable = true)
-- review_count: long (nullable = true)
-- stars: double (nullable = true)
-- state: string (nullable = true)

```

Display the first 5 rows with the following columns:

- business_id
- name
- city
- state
- categories

```
In [6]: busi_df = business_df.select('business_id', 'name', 'city', 'state', 'stars', 'categories')
busi_df.show(5)
```

```

+-----+-----+-----+-----+-----+-----+
---+
|      business_id|      name|      city|state|stars|      categories|
+-----+-----+-----+-----+-----+-----+
---+
|f9NumwFMBDn751xgF...|The Range At Lake...|    Cornelius|    NC|   3.5|Active Life, Gu
n/...|
|YzvJg0SayhoZgCljU...|  Carlos Santo, NMD|    Scottsdale|    AZ|   5.0|Health & Medica

```

```

1,...|
|XNoUzKckATkOD1hP6...|          Felinus|          Montreal|    QC|    5.0|Pets, Pet Servic
e...|
|60AZjbxqM5o129BuH...|Nevada House of Hose|North Las Vegas|    NV|    2.5|Hardware Stores,
...|
|51M2Kk903DFYI6gnB...|USE MY GUY SERVIC...|          Mesa|    AZ|    4.5|Home Services, P
1...|
+-----+-----+-----+-----+-----+-----+
---+
only showing top 5 rows

```

Part II: Analyzing Categories

Let's now answer: **How many unique categories are represented in this dataset?**

Essentially, we have the categories per business as a list - this is useful to quickly see what each business might be represented as but it is difficult to easily answer the following questions such as:

- How many businesses are categorized as Active Life?
- What are the top 20 most popular categories available?

1. Association Table

We need to "break out" these categories from the business ids? One common approach to take is to build an association table mapping a single business id multiple times to each distinct category.

For instance, given the following:

business_id	categories
abcd123	a,b,c

We would like to derive something like:

business_id	category
abcd123	a
abcd123	b
abcd123	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Display the first 5 rows of the association table below

```
In [7]: associ_table_one = business_df.select('business_id', explode(split(business_df.categories)))
associ_table_one.show(5)
```

```

+-----+-----+
|          business_id|          category|
+-----+-----+
|f9NumwFMBDn751xgF...|    Active Life|
|f9NumwFMBDn751xgF...|Gun/Rifle Ranges|
|f9NumwFMBDn751xgF...|    Guns & Ammo|

```

```
|f9NumwFMBDn751xgF...| Shopping|
|YzvJg0SayhoZgCljU...| Health & Medical|
+-----+
only showing top 5 rows
```

2. Total Unique Categories

Finally, we are ready to answer the question: **what is the total number of unique categories available?**

```
In [8]: associ_table_one.select('category').distinct().count()
```

1336

3. Top Categories By Business

Now let's find the top categories in this dataset by rolling up categories.

Counts of Businesses / Category

```
In [9]: category_count = associ_table_one.select('category').groupby(associ_table_one.category)
category_count.show()
```

```
+-----+-----+
|      category|count|
+-----+-----+
| Paddleboarding|   36|
| Dermatologists|  341|
| Aerial Tours  |   28|
| Hobby Shops   |  828|
| Bubble Tea    |  720|
| Embassy       |   13|
| Tanning       |  938|
| Handyman      |  682|
| Aerial Fitness|   29|
| Falafel       |  159|
| Outlet Stores |  399|
| Summer Camps |  318|
| Clothing Rental|   55|
| Sporting Goods| 2311|
| Cooking Schools|  118|
| College Counseling|  15|
| Lactation Services|  50|
| Ski & Snowboard S...|  50|
| Museums       |  359|
| Doulas        |   45|
+-----+-----+
only showing top 20 rows
```

Bar Chart of Top Categories


With this data available, let us now build a barchart of the top 20 categories

```
In [52]: top_20 = category_count.sort(desc('count')).limit(20).toPandas()

ax = top_20.plot(kind='barh', x='category', y='count',
                  figsize=(8, 4), color='#5fa8d3', zorder=2, width=0.85)
```

```
ax.invert_yaxis()
ax.set_xlabel("Count")
ax.set_ylabel("Category")
ax.set_title("Top 20 Categories By Business")
ax.get_legend().remove()

plt.tight_layout()
%matplotlib plt
```



Part III. Do Yelp Reviews Skew Negative?

Oftentimes, it is said that the only people who write a written review are those who are extremely dissatisfied or extremely satisfied with the service received.

How true is this really? Let's try and answer this question.

1. Loading Review Data

Begin by loading the review data set from S3 and printing schema to determine what data is available

```
In [11]: review_df = spark.read.json('s3://yelpreviewdataset/yelp_academic_dataset_review.json')
review_df.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

Let's begin by listing the business_id and stars columns together for the user reviews data

```
In [12]: business_stars = review_df.select('business_id', 'stars')
business_stars.show(5)
```

```
+-----+-----+
|      business_id|stars|
+-----+-----+
|-MhFebM0QIsKt87iD...| 2.0|
|lbrU8StCq3yDfr-QM...| 1.0|
|HQL28KMwrEKHqhFrr...| 5.0|
|5JxlZaqCnk1MnbgRi...| 1.0|
|IS4cv902ykd8wj1TR...| 4.0|
+-----+-----+
only showing top 5 rows
```

Now, let's aggregate along the stars column to get a resultant dataframe that displays average stars per business as accumulated by users who **took the time to submit a written review**

```
In [13]: written_review = review_df.where(col("text").isNotNull()).groupby(review_df.business_id)
written_review.show(5)
```

```
+-----+-----+
|      business_id|      avg(stars)|
+-----+-----+
|RMjCnixEY5i12Ciqn...| 3.5316455696202533|
|VHsNB3pdGVcRgs6C3...| 3.411764705882353|
|kpbhERZoj1eTDRnMV...| 2.0333333333333333|
|ipFreSFhjClfNETuM...| 2.6|
|9A_mB7Ez3RIh26EN5...| 2.6|
+-----+-----+
only showing top 5 rows
```

Now the fun part - let's join our two dataframes (reviews and business data) by business_id

```
In [14]: user_review = review_df.groupby(review_df.business_id).agg(avg(col("stars")))
new_df1 = business_df.join(written_review, on=['business_id'])
new_df2= business_df.join(user_review,on=['business_id'])
new_df1.select("""avg(stars)""", "stars", "name", "city", "state").sort(desc("""avg(stars)"""))
```

```
+-----+-----+-----+-----+-----+
|avg(stars)|stars|      name|      city|state|
+-----+-----+-----+-----+-----+
|      5.0| 5.0| Larry Fafalak, LMT| Las Vegas| NV|
|      5.0| 5.0| Rentech Solutions| Willoughby| OH|
|      5.0| 5.0| Everest Curry and...| Calgary| AB|
|      5.0| 5.0| Krown Rust Contro...| Markham| ON|
|      5.0| 5.0| EVO Swim School| Mesa| AZ|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Compute a new dataframe that calculates what we will call the skew (for lack of a better word) between the avg stars accumulated from written reviews and the actual star rating of a business (ie: the average of stars given by reviewers who wrote an actual review and reviewers who just provided a star rating).

The formula you can use is something like:


```
(row['avg(stars)'] - row['stars']) / row['stars']
```


If the **skew** is negative, we can interpret that to be: reviewers who left a written response were more dissatisfied than normal. If **skew** is positive, we can interpret that to be: reviewers who left a written response were more satisfied than normal.

```
In [15]: fv_df1 = new_df1.select("avg(stars)", "stars", "name", "city", "state").sort("avg(stars)")
fv_df2 = new_df2.select("avg(stars)", "stars", "name", "city", "state").sort("avg(stars)")
fv_df = fv_df1.withColumn("skew", ((fv_df1["avg(stars)"] - fv_df2["stars"]) / fv_df1["s
```

And finally, graph it!

```
In [50]: plt.figure(figsize=(8,4))
ax = sns.distplot(fv_df["skew"])
ax.set_xlabel('Skew')
plt.title("Yelp Review Skew")

%matplotlib plt
```



So, do Yelp (written) Reviews skew negative? Does this analysis actually prove anything? Expound on implications / interpretations of this graph.

The distribution of skew appears to be normal, but skewed a little bit to the right. The implications of the above graph are that the satisfaction level of reviewers who left positively skewed reviews is greater than the dissatisfaction level of reviewers who left negatively skewed reviews.

Part IV. Should the Elite be Trusted?

How accurate or close are the ratings of an "elite" user (check Users table schema) vs the actual business rating

It takes a special Yelper to become an Elite. Frequent, quality reviews and photos are important in the application of the elite status on Yelp. Elite candidates need to meet the criteria below for the consideration.

To become Elite, Yelpers agree that they

- Are using their real name on Yelp.
- Have a clear photo of themselves on their profile page.
- Are of legal drinking age where they live.

They also agree that they are NOT

- A business owner.
- Closely affiliated with a business owner.
- Managing a Yelp Business Account.
- Working for one of Yelp's competitors.

It's important to know that accepting compensation or freebies in exchange for reviews or leveraging the Elite Squad for personal or commercial gain will result in Elite status being revoked or account closure.

1. Loading User Data

```
In [17]: user_df = spark.read.json('s3://yelpreviewdataset/yelp_academic_dataset_user.json')
```

2. Overview of Data

```
In [18]: user_df.printSchema()
print(f'User Dataset Columns: {len(user_df.columns)} | Rows: {user_df.count():,}')
review_df.printSchema()
print(f'Review Dataset Columns: {len(review_df.columns)} | Rows: {review_df.count():,}')
```

```
root
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
```

```
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)
```

User Dataset Columns: 22 | Rows: 1,968,703

root

```
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

Review Dataset Columns: 9 | Rows: 8,021,122

```
In [19]: user_df.select('user_id', 'elite').show(5)
```

```
+-----+-----+
|          user_id|          elite|
+-----+-----+
| ntlvfPzc8eglqv92...|          |
| FOBRP1BHa3WPHFB5q...| 2008,2009,2010,20...|
| zZUnPeh2hEp0WydbA...|          2010|
| QaELAmRcDc5TfJEyl...|          2009|
| xvu8G900tezTzbbfq...| 2009,2010,2011,20...|
+-----+-----+
```

only showing top 5 rows

3. Split Elite column

```
In [20]: user_elite_split = user_df.select('user_id', explode(split(user_df.elite, ',')).alias('
user_elite_split = user_elite_split.withColumn('elite', user_elite_split.elite.cast(Int
user_elite_split.show(5)
print(f'User Elite Split Dataset Columns: {len(user_elite_split.columns)} | Rows: {user
```

```
+-----+-----+
|          user_id|elite|
+-----+-----+
| ntlvfPzc8eglqv92...| null|
| FOBRP1BHa3WPHFB5q...| 2008|
| FOBRP1BHa3WPHFB5q...| 2009|
| FOBRP1BHa3WPHFB5q...| 2010|
| FOBRP1BHa3WPHFB5q...| 2011|
+-----+-----+
```

only showing top 5 rows

User Elite Split Dataset Columns: 2 | Rows: 2,125,315

```
In [21]: user_elite_split.select("elite").distinct().sort('elite', ascending=False).show()
```

```
+-----+
|elite|
+-----+
| 2018|
| 2017|
| 2016|
```

```
| 2015|
| 2014|
| 2013|
| 2012|
| 2011|
| 2010|
| 2009|
| 2008|
| 2007|
| 2006|
| null|
+-----+
```

```
In [22]: Elite_or_Not = user_elite_split.select('user_id',
        when( user_elite_split.elite.isNull(), "Not Elite").otherwise("Elite").alias
        Elite_or_Not.show()
```

```
+-----+-----+
|          user_id|Elite or Not|
+-----+-----+
|ntlvfPzc8eglqvk92...|    Not Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|FOBRP1BHa3WPHFB5q...|      Elite|
|zZUnPeh2hEp0WydbA...|      Elite|
|QaELAmRcDc5TfJEy1...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|xvu8G900tezTzbbfq...|      Elite|
|z5_82komKV3mI4ASG...|      Elite|
|ttumcu6hWshk_EJWV...|    Not Elite|
+-----+-----+
only showing top 20 rows
```

```
In [23]: unique_user_df = Elite_or_Not.dropDuplicates(['user_id'])
        unique_user_df.show()
```

```
+-----+-----+
|          user_id|Elite or Not|
+-----+-----+
|---RfKzBwQ8t3wu-L...|    Not Elite|
|--1UpCuUDJQbqiuFX...|    Not Elite|
|--AGAPpP1pgp1afbq...|    Not Elite|
|--C-42rr7hPSsUR0J...|    Not Elite|
|--ChzqcPs4YFWlw1j...|    Not Elite|
|--ET3paBtrThD95dk...|    Not Elite|
|--GLTFzU93A40YB56...|    Not Elite|
|--I4wRDhmM2J2VLzK...|    Not Elite|
|--RquisWmBzcezXZr...|    Not Elite|
|--UizzbnQlZg7bEv2...|    Not Elite|
|--cd_gA-9Q8gM9P2c...|    Not Elite|
|--dhSVoOFDBiMCCwD...|    Not Elite|
|--fpTdHQOGWGbAjk9...|    Not Elite|
```

```
--ju6XpRd0dY1Swmf...| Not Elite|
--oVdTxD7QVr8Y0U...| Not Elite|
--pWqE-KOWDwo5ADG...| Not Elite|
--t6W1JHbStaCp5RO...| Not Elite|
--tmwndDOZJwFRvvt...| Not Elite|
--yrdC1dIR6VYsW6k...| Not Elite|
-06viLTmtlRTHxxDg...| Not Elite|
+-----+
only showing top 20 rows
```

4. Join "Unique User" Dataset with Review Dataset

```
In [24]: user_join_review = review_df.join(unique_user_df, on = "user_id", how='left')
print(f'User Join Review Dataset Columns: {len(user_join_review.columns)} | Rows: {user
```

User Join Review Dataset Columns: 10 | Rows: 8,021,122

5. Clean Data

Combined Datasets which includes elite and non-elite

```
In [25]: combine_df = user_join_review.select('review_id', 'business_id', 'stars', 'user_id', 'Elite')
combine_df.show()
```

```
+-----+-----+-----+-----+
| review_id | business_id | stars | user_id | Elite or Not |
+-----+-----+-----+-----+
| rv2EaVEP_cs0Yzc-z... | Z3ZSar8IVAR2qIupq... | 5.0 | ---RfKzBwQ8t3wu-L... | Not Elite |
| HVR4EWzZMlyPrdbzE... | kJhQq1BFz7lOYLve7... | 1.0 | --1UpCuUDJQbqiuFX... | Not Elite |
| uy83M2YEnInksqsKX... | EpPOZAG0u7qHP-jv5... | 5.0 | --1UpCuUDJQbqiuFX... | Not Elite |
| EHsBHPADGfll02Zm5... | OLmcIJ7VBCxaYhZSN... | 5.0 | --AGAppP1pgp1afbq... | Not Elite |
| xtHcnwOx-27sunclu... | WoiOpMEcbAf0qNYXq... | 5.0 | --AGAppP1pgp1afbq... | Not Elite |
| pFq8ijDeB-Gz1HXsS... | L--9JNAb6UDyq7wa... | 4.0 | --C-42rr7hPSsUROJ... | Not Elite |
| SI_ONkbwzN_i38Gvg... | 4KmrhtfnngTVFa2d... | 4.0 | --ChzqcPs4YFWlw1j... | Not Elite |
| fHqAyF58eC6vC4_BP... | AMTNJbYbu0OMMAkx4... | 4.0 | --ChzqcPs4YFWlw1j... | Not Elite |
| bQkvjklZmtFYaYd0... | KVs8wRgnLX8QWoNZ... | 3.0 | --ChzqcPs4YFWlw1j... | Not Elite |
| V4nVpftxljW4sF0g0... | 6pG7n8Rx_7ZxeQQk6... | 2.0 | --ChzqcPs4YFWlw1j... | Not Elite |
| YSW-S2XUyCKR3jUtW... | F9CcIFltPDXiOkCCF... | 4.0 | --ChzqcPs4YFWlw1j... | Not Elite |
| mfqVYzvoeiZREW8bs... | QZV9hW3WP9o9SmmV2... | 5.0 | --ET3paBtrThD95dk... | Not Elite |
| 99Vpr7r8dGR0txvL3... | pT6baSMzC6rZfwhp... | 5.0 | --GLTFzU93A40YB56... | Not Elite |
| YQN6mfSAX12LFsn6r... | JmI9nslLD7KZqRr... | 2.0 | --I4wRDhmM2J2VLzK... | Not Elite |
| X2sbxAYTM9KYjyP0e... | HW7JPZBImm3tyEpDg... | 5.0 | --RquisWmBzcezXZr... | Not Elite |
| cqrmoHebDTzgc5hj0... | XNFA-aJFX8IQjol8D... | 4.0 | --RquisWmBzcezXZr... | Not Elite |
| ubpg7b5NJUiH_A_2d... | W2Vis19kUa7kP6GkS... | 5.0 | --RquisWmBzcezXZr... | Not Elite |
| Bz_KEvFEyKL1QtBFe... | hDD6-yk1yuuRiVfdt... | 2.0 | --UizzbnQ1Zg7bEv2... | Not Elite |
| sZR9FQeM1c07UKhTD... | eNFubUPJR7yIQah-N... | 4.0 | --cd_gA-9Q8gM9P2c... | Not Elite |
| yhgRUG0ctQ0aEaaIi... | uPa5hrWmHm0n114MS... | 4.0 | --cd_gA-9Q8gM9P2c... | Not Elite |
+-----+-----+-----+-----+
only showing top 20 rows
```

Combined (Elite and Non-Elite) Average Ratings Grouped by Business ID

```
In [26]: combine_stars_df = combine_df.groupBy("business_id").agg(F.mean('stars').alias('Stars'))
combine_stars_df.show()
```

```
+-----+-----+
| business_id | Stars |
+-----+-----+
| --9e10NYQuAa-CB_R... | 4.11784140969163 |
```

```

RtUvSWO_UZ8V3Wpj0... 4.133498145859085
kpbhERZoj1eTDRnMV... 2.0333333333333333
eKznX8VTfcQrjCqXp... 4.3584905660377355
umwULmdsxx8aTsoRQ... 2.388888888888889
ru_WUOAmx9xPBxcJu... 5.0
VHsNB3pdGVcRgs6C3... 3.411764705882353
rtwojGcYuhbLbQ9D1... 3.4545454545454546
SjgeuBlgKER9yegpo... 3.973643410852713
VmSrPPO2WXmOKjUW7... 3.227906976744186
0FWYa5RT_gQOwW3CR... 3.4545454545454546
RMjCnixEY5i12Ciqn... 3.5316455696202533
l1CxryWr8j1S39tus... 4.43839541547278
35X1ZV9tSEqB__yJE... 3.0316742081447963
X6jKcN5FoRiJ1t7y4... 1.8444444444444446
DgCAM01n2Qo5DsoKj... 3.3448275862068964
xusE_x84QOEDaRZ8r... 3.7096774193548385
cz5vz-893D3LNNH3TM... 3.803514376996805
jfdUtdkXogP2kjK5K... 3.6323529411764706
uC3qwaxsOkdJzp0c0... 3.368948247078464

```

```

+-----+
only showing top 20 rows

```

Elite Only Dataset

```

In [27]: elite_df = combine_df.filter(col("Elite or Not") == "Elite")
         elite_df.show()

```

```

+-----+-----+-----+-----+-----+
| review_id | business_id | stars | user_id | Elite or Not |
+-----+-----+-----+-----+-----+
| EIKPUavToyh-dz2eE... | WYw3Uf56DT5IwpaLN... | 5.0 | -1_RJoRLeoDK3h_gN... | Elite |
| ygfB-2RWSKtI3jVC3... | 0gXYLVPNWz0WT8wXQ... | 4.0 | -1_RJoRLeoDK3h_gN... | Elite |
| 84GE9SrQCw-Yv-qpM... | W2CzAePJakvARgoQu... | 3.0 | -1_RJoRLeoDK3h_gN... | Elite |
| tyTkxTaNH1sL8t9XK... | iCQpiavjjPzJ5_3gP... | 4.0 | -1_RJoRLeoDK3h_gN... | Elite |
| TJDpUewi8F1E9eUgi... | qalkZ4AQDWzYrFvQV... | 5.0 | -1_RJoRLeoDK3h_gN... | Elite |
| bAd_-cPcZNsVfhFgN... | _w5hBpkjHs5_Hv3pL... | 4.0 | -1_RJoRLeoDK3h_gN... | Elite |
| kKuzCM7kpGqCue3iD... | Yl05MqCs9xRzrJfK... | 5.0 | -1_RJoRLeoDK3h_gN... | Elite |
| ITIUKGvnRE3u6RLns... | 7FvDsYqtij_BbaGVt... | 3.0 | -1_RJoRLeoDK3h_gN... | Elite |
| W4FCaD23_CzAoC28j... | A4zLP5AyKEEHQr_dw... | 4.0 | -1_RJoRLeoDK3h_gN... | Elite |
| 6aNCf2uoLiLz27pWS... | 90bL34o2KEes9pUnC... | 4.0 | -1_RJoRLeoDK3h_gN... | Elite |
| 3QvS6Ued-M_5wjln... | fE9SP84G6TZrv36FL... | 3.0 | -1_RJoRLeoDK3h_gN... | Elite |
| yUWEX8m3DnwI3YnNW... | MBekdd_f7S1ezEzZb... | 5.0 | -1xh431AhmrByuMzc... | Elite |
| 23fDyVgPz7-gHvNvx... | deL9fV4Jw3XhS0WqG... | 4.0 | -1xh431AhmrByuMzc... | Elite |
| my4UdVCRQ9dITsWRO... | mz9ltimeAIy2c2qf5... | 5.0 | -1xh431AhmrByuMzc... | Elite |
| 0fWB1f-2BK9fmgYTA... | M4vh_kzppP1nsxo7h... | 3.0 | -1xh431AhmrByuMzc... | Elite |
| wF-_nw2kG_vQ0079N... | deL9fV4Jw3XhS0WqG... | 4.0 | -1xh431AhmrByuMzc... | Elite |
| gJeVSSm1CQ6XOLh0v... | KdQM64AQ5_ppgs6Ro... | 4.0 | -1xh431AhmrByuMzc... | Elite |
| X_JpVPD3EoPF8YRpb... | LYNNKn14jAiU1-U-9... | 4.0 | -1xh431AhmrByuMzc... | Elite |
| PKTOJZZvFpKaKmitR... | Bkkwt8E9MHvgCHn41... | 5.0 | -1xh431AhmrByuMzc... | Elite |
| qIhEdr18_bLGuaiRL... | TqUVH70x_3qEkCxCC... | 4.0 | -1xh431AhmrByuMzc... | Elite |

```

```

+-----+
only showing top 20 rows

```

Elite Average Rating Grouped by Business ID

```

In [28]: elite_stars_df = elite_df.groupBy("business_id").agg(F.mean('stars').alias('Stars rated
         elite_stars_df.show()

```

```

+-----+-----+
| business_id | Stars rated by elite |
+-----+-----+
| eKznX8VTfcQrjCqXp... | 4.268817204301075 |

```

RtUvSWO_UZ8V3Wpj0...	4.156193895870736
rtwojGcYuhbLbQ9D1...	3.3636363636363638
--9e10NYQuAa-CB_R...	4.1916058394160585
X6jKcN5FoRiJ1t7y4...	4.5
SjgeuBlgKER9yegpo...	3.8938775510204082
jfdUtdkXogP2kjK5K...	3.3846153846153846
uC3qwaxsOkdJzp0c0...	3.6745562130177514
yJGr280XuMk2bCKY1...	3.125
f4mh1Y0rnnvbjRfQ3j...	3.875
cz5vz-893D3LNH3TM...	3.8587570621468927
MEoDTsA3Af6TLzB7Z...	3.2142857142857144
qtsrM6Xxh1LqxG0X6...	4.5
0juzFQpprqmuapKh6...	3.6
VmSrPP02WxmOKjUW7...	3.423076923076923
Rxb7oKtKyDUwFNc2...	3.4285714285714284
4iY_gyKX2ogbem7ra...	4.4444444444444445
mx0Pjm0erpv1CqsRI...	3.8
VHsNB3pdGVcRgs6C3...	4.0
LCRdP3m826-Df52-x...	1.0

only showing top 20 rows

Non-Elite Dataset

```
In [29]: non_elite_df = combine_df.filter(col("Elite or Not") == "Not Elite")
non_elite_df.show()
```

review_id	business_id	stars	user_id	Elite or Not
rv2EaVEP_cs0Yzc-z...	Z3ZSar8IVAR2qIupq...	5.0	---RfKzBwQ8t3wu-L...	Not Elite
HVR4EWzZMlyPrdbzE...	kJhQq1BFz7l0YLve7...	1.0	--1UpCuUDJQbqiuFX...	Not Elite
uy83M2YEnInksqsKX...	EpPOZAG0u7qHP-jv5...	5.0	--1UpCuUDJQbqiuFX...	Not Elite
EHsBHPADGfll02Zm5...	OLmcIJ7VBCxaYhZSN...	5.0	--AGAppP1pgp1afbq...	Not Elite
xtHcnw0x-27sunclu...	WoiOpMEcbAf0qNYXq...	5.0	--AGAppP1pgp1afbq...	Not Elite
pFq8ijDeB-Gz1HXsS...	L--9JNAb6UDyq7wa...	4.0	--C-42rr7hPSsUROJ...	Not Elite
V4nVpftxljW4sF0g0...	6pG7n8Rx_7ZXeQQk6...	2.0	--ChzqcPs4YFWlw1j...	Not Elite
fHqAyF58eC6vC4_BP...	AMTNJbYbu0OMMAkx4...	4.0	--ChzqcPs4YFWlw1j...	Not Elite
YSW-S2XUyCKR3jUtW...	F9CcIFltPDXi0kCCF...	4.0	--ChzqcPs4YFWlw1j...	Not Elite
SI_ONkbnwN_i38Gvg...	4KmrtrhtfnngTVFa2d...	4.0	--ChzqcPs4YFWlw1j...	Not Elite
bQkvjklZmtFYaYd0...	KVsv8wRGnLX8QWoNz...	3.0	--ChzqcPs4YFWlw1j...	Not Elite
mfqVYzvoeiZREW8bs...	QZV9hW3WP9o9SmmV2...	5.0	--ET3paBtrThD95dk...	Not Elite
99Vpr7r8dGR0txvL3...	pT6baSMzC6rZfwhp...	5.0	--GLTFzU93A40YB56...	Not Elite
YQN6mfSAX12LFsn6r...	JmI9nslLD7KZqRr...	2.0	--I4wRDhmM2J2VLzK...	Not Elite
X2sbxAYTM9KYjyP0e...	HW7JPZBImm3tyEpDg...	5.0	--RquisWmBzcezXZr...	Not Elite
cqrmoHebDTzgc5hj0...	XNFA-aJFX8IQjol8D...	4.0	--RquisWmBzcezXZr...	Not Elite
ubpg7b5NJUiH_A_2d...	W2Vis19kUa7kP6GkS...	5.0	--RquisWmBzcezXZr...	Not Elite
Bz_KEvFEyKL1QtbfFe...	hDD6-yk1yuuRIvfdt...	2.0	--UizzbnQ1Zg7bEv2...	Not Elite
PR01xlQ0srxmQ8TIu...	9Eghhu_LzEJgDKNgi...	4.0	--cd_gA-9Q8gM9P2c...	Not Elite
Ct03r0f40jz05T1jm...	fQwB9Z98YEhkJit7c...	3.0	--cd_gA-9Q8gM9P2c...	Not Elite

only showing top 20 rows

Non-Elite Average Rating Grouped by Business ID

```
In [30]: non_elite_stars_df = non_elite_df.groupBy("business_id").agg(F.mean('stars').alias('Sta
non_elite_stars_df.show()
```

business_id	Stars rated by non elite
--9e10NYQuAa-CB_R...	4.08596214511041


RtUvSWO_UZ8V3Wpj0...	4.121583411875589
kpbhERZoj1eTDRnMV...	1.9642857142857142
umwULmdsxx8aTsoRQ...	2.25
ru_WUOAmx9xPBxcJu...	5.0
VHsNB3pdGVcRgs6C3...	3.279279279279279
eKznX8VTfcQrjCqXp...	4.406976744186046
SjgeuBlgKER9yegpo...	4.0225
VmSrPPO2WxmOKjUW7...	3.201058201058201
0FWYa5RT_gQOwW3CR...	3.4
RMjCnixEY5i12Ciqn...	3.6226415094339623
l1CxryWr8j1S39tus...	4.4627831715210355
35X1ZV9tSEqB_yJE...	3.0080645161290325
X6jKCn5FoRiJ1t7y4...	1.7209302325581395
DgCAM01n2Qo5DsoKj...	3.1739130434782608
xusE_x84QOEaRZ8r...	3.7142857142857144
cz5vz-893D3LNH3TM...	3.7817371937639197
uC3qwaxsOkdJzpOc0...	3.2488372093023257
rtwojGcYuhbLbQ9D1...	3.6363636363636362
D-CyKxLr_NJe1fI4e...	1.8235294117647058

+-----+
only showing top 20 rows

Prepare data for plotting

```
In [31]: combined_data = combine_stars_df.toPandas()["Stars"].values.tolist()
elite_data = elite_stars_df.toPandas()["Stars rated by elite"].values.tolist()
non_elite_data = non_elite_stars_df.toPandas()["Stars rated by non elite"].values.tolis
data = [combined_data, elite_data, non_elite_data]
```

```
In [49]: fig = plt.figure(figsize =(8, 4))
plt.boxplot(data)
plt.xticks([1, 2, 3], ['Combined Rating', 'Elite Rating', 'Non-Elite Rating'])
plt.title("Comparison of Combined Rating, Elite Rating and Non-Elite Rating")
y_ticks = np.arange(0, 6, 1)
plt.yticks(y_ticks)
plt.show()
%matplotlib plt
```



As we can see from the above boxplot, elite data has more outliers. Additionally, the first, third quantiles and the median of the elite ratings are also higher than the non-elites' ratings. From my point of view, I would say elite should not be trusted.